

# Business Problem

## Description

Quora is a place to gain and share knowledge—about anything. It's a platform to ask questions and connect with people who contribute unique insights and quality answers. This empowers people to learn from each other and to better understand the world.

Over 100 million people visit Quora every month, so it's no surprise that many people ask similarly worded questions. Multiple questions with the same intent can cause seekers to spend more time finding the best answer to their question, and make writers feel they need to answer multiple versions of the same question. Quora values canonical questions because they provide a better experience to active seekers and writers, and offer more value to both of these groups in the long term.

Credits: Kaggle

### Problem Statement

- Identify which questions asked on Quora are duplicates of questions that have already been asked.
- This could be useful to instantly provide answers to questions that have already been answered.
- We are tasked with predicting whether a pair of questions are duplicates or not.

## Sources/Useful Links

- Source : <https://www.kaggle.com/c/quora-question-pairs> (<https://www.kaggle.com/c/quora-question-pairs>)
- **Useful Links**
- Discussions : <https://www.kaggle.com/anokas/data-analysis-xgboost-starter-0-35460-lb/comments> (<https://www.kaggle.com/anokas/data-analysis-xgboost-starter-0-35460-lb/comments>)
- Kaggle Winning Solution and other approaches: <https://www.dropbox.com/sh/93968nfnrzh8bp5/AACZdtsApc1QSTQc7X0H3QZ5a?dl=0> (<https://www.dropbox.com/sh/93968nfnrzh8bp5/AACZdtsApc1QSTQc7X0H3QZ5a?dl=0>)
- Blog 1 : <https://engineering.quora.com/Semantic-Question-Matching-with-Deep-Learning> (<https://engineering.quora.com/Semantic-Question-Matching-with-Deep-Learning>)
- Blog 2 : <https://towardsdatascience.com/identifying-duplicate-questions-on-quora-top-12-on-kaggle-4c1cf93f1c30> (<https://towardsdatascience.com/identifying-duplicate-questions-on-quora-top-12-on-kaggle-4c1cf93f1c30>)

## Real world/Business Objectives and Constraints

1. The cost of a mis-classification can be very high.
2. You would want a probability of a pair of questions to be duplicates so that you can choose any threshold of choice.
3. No strict latency concerns.
4. Interpretability is partially important.

# Machine Learning Problm

## Data

### Data Overview

- Data will be in a file Train.csv
- Train.csv contains 5 columns : qid1, qid2, question1, question2, is\_duplicate
- Size of Train.csv - 60MB
- Number of rows in Train.csv = 404,290

### Example Data point

```
"id","qid1","qid2","question1","question2","is_duplicate"
"0","1","2","What is the step by step guide to invest in share market in india?","What is the step by step guide to invest in share market?","0"
"1","3","4","What is the story of Kohinoor (Koh-i-Noor) Diamond?","What would happen if the Indian government stole the Kohinoor (Koh-i-Noor) diamond back?","0"
"7","15","16","How can I be a good geologist?","What should I do to be a great geologist?","1"
"11","23","24","How do I read and find my YouTube comments?","How can I see all my Youtube comments?","1"
```

## Mapping the real world problem to an ML problem

### Type of Machine Leaning Problem

It is a binary classification problem, for a given pair of questions we need to predict if they are duplicate or not.

### Performance Metric

Source: <https://www.kaggle.com/c/quora-question-pairs#evaluation> (<https://www.kaggle.com/c/quora-question-pairs#evaluation>)

Metric(s):

- log-loss : <https://www.kaggle.com/wiki/LogarithmicLoss> (<https://www.kaggle.com/wiki/LogarithmicLoss>)
- Binary Confusion Matrix

## Exploratory Data Analysis

```
In [0]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from subprocess import check_output
%matplotlib inline
import plotly.offline as py
py.init_notebook_mode(connected=True)
import plotly.graph_objs as go
import plotly.tools as tls
import os
import gc

import re
from nltk.corpus import stopwords
import distance
from nltk.stem import PorterStemmer
from bs4 import BeautifulSoup
```

## Reading data and basic stats

```
In [0]: df = pd.read_csv("train.csv")

print("Number of data points:",df.shape[0])
```

Number of data points: 404290

```
In [0]: df.head()
```

Out[8]:

	id	qid1	qid2	question1	question2	is_duplicate
0	0	1	2	What is the step by step guide to invest in sh...	What is the step by step guide to invest in sh...	0
1	1	3	4	What is the story of Kohinoor (Koh-i-Noor) Dia...	What would happen if the Indian government sto...	0
2	2	5	6	How can I increase the speed of my internet co...	How can Internet speed be increased by hacking...	0
3	3	7	8	Why am I mentally very lonely? How can I solve...	Find the remainder when $23^{24}$ is di...	0
4	4	9	10	Which one dissolve in water quikly sugar, salt...	Which fish would survive in salt water?	0

```
In [0]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 404290 entries, 0 to 404289
Data columns (total 6 columns):
id                404290 non-null int64
qid1              404290 non-null int64
qid2              404290 non-null int64
question1         404290 non-null object
question2         404288 non-null object
is_duplicate      404290 non-null int64
dtypes: int64(4), object(2)
memory usage: 18.5+ MB
```

We are given a minimal number of data fields here, consisting of:


- id: Looks like a simple rowID
- qid{1, 2}: The unique ID of each question in the pair
- question{1, 2}: The actual textual contents of the questions.
- is\_duplicate: The label that we are trying to predict - whether the two questions are duplicates of each other.

## Distribution of data points among output classes

- Number of duplicate(smilar) and non-duplicate(non similar) questions

```
In [0]: df.groupby("is_duplicate")["id"].count().plot.bar()
```

Out[10]: <matplotlib.axes.\_subplots.AxesSubplot at 0x22b00727d30>



```
In [0]: print("Number of question pairs available for training is",len(df))
```

Number of question pairs available for training is 404290

```
In [0]: print("Percentage of question pairs that are not similar (is_duplicate = 0) is",100 - round(df['is_duplicate'].mean()*100, 2))
print("Percentage of question pairs that are similar (is_duplicate = 1) is",round(df['is_duplicate'].mean()*100, 2))
```

Percentage of question pairs that are not similar (is\_duplicate = 0) is 63.08%  
Percentage of question pairs that are similar (is\_duplicate = 1) is 36.92%

## Number of unique questions

```
In [0]: qids = pd.Series(df['qid1'].tolist() + df['qid2'].tolist())
unique_qs = len(np.unique(qids))
qs_morethan_onetime = np.sum(qids.value_counts() > 1)
print ('Total number of Unique Questions are: {}'.format(unique_qs))
#print len(np.unique(qids))

print ('Number of unique questions that appear more than one time: {} ({}%)'.format(qs_morethan_onetime,qs_morethan_onetime/unique_qs*100))

print ('Max number of times a single question is repeated: {}'.format(max(qids.value_counts()))))

q_vals=qids.value_counts()

q_vals=q_vals.values


Total num of Unique Questions are: 537933

Number of unique questions that appear more than one time: 111780 (20.77953945937505%)

Max number of times a single question is repeated: 157
```

```
In [0]: x = ["unique_questions" , "Repeated Questions"]
y = [unique_qs , qs_morethan_onetime]

plt.figure(figsize=(10, 6))
plt.title ("Plot representing unique and repeated questions ")
sns.barplot(x,y)
plt.show()
```



Checking for Duplicates

```
In [0]: #checking whether there are any repeated pair of questions

pair_duplicates = df[['qid1','qid2','is_duplicate']].groupby(['qid1','qid2']).count().reset_index()

print ("Number of duplicate questions",(pair_duplicates).shape[0] - df.shape[0])

Number of duplicate questions 0
```

Number of occurrences of each question

```
In [0]: plt.figure(figsize=(20, 10))

plt.hist(qids.value_counts(), bins=160)

plt.yscale('log', nonposy='clip')


plt.title('Log-Histogram of question appearance counts')

plt.xlabel('Number of occurences of question')

plt.ylabel('Number of questions')

print ('Maximum number of times a single question is repeated: {}'.format(max(qids.value_counts()))))

Maximum number of times a single question is repeated: 157
```



Checking for NULL values

In [0]:

```
#Checking whether there are any rows with null values
nan_rows = df[df.isnull().any(1)]
print (nan_rows)
```

	id	qid1	qid2	question1	question2	\
105780	105780	174363	174364	How can I develop android app?	NaN	
201841	201841	303951	174364	How can I create an Android app?	NaN	

	is_duplicate
105780	0
201841	0

- There are two rows with null values in question2

In [0]:

```
# Filling the null values with ' '
df = df.fillna(' ')
nan_rows = df[df.isnull().any(1)]
print (nan_rows)
```

Empty DataFrame  
Columns: [id, qid1, qid2, question1, question2, is\_duplicate]  
Index: []

## Basic Feature Extraction (before cleaning)

Let us now construct a few features like:

- freq\_qid1** = Frequency of qid1's
- freq\_qid2** = Frequency of qid2's
- q1len** = Length of q1
- q2len** = Length of q2
- q1\_n\_words** = Number of words in Question 1
- q2\_n\_words** = Number of words in Question 2
- word\_Common** = (Number of common unique words in Question 1 and Question 2)
- word\_Total** =(Total num of words in Question 1 + Total num of words in Question 2)
- word\_share** = (word\_common)/(word\_Total)
- freq\_q1+freq\_q2** = sum total of frequency of qid1 and qid2
- freq\_q1-freq\_q2** = absolute difference of frequency of qid1 and qid2

In [0]:

```
if os.path.isfile('df_fe_without_preprocessing_train.csv'):
    df = pd.read_csv("df_fe_without_preprocessing_train.csv",encoding='latin-1')
else:
    df['freq_qid1'] = df.groupby('qid1')['qid1'].transform('count')
    df['freq_qid2'] = df.groupby('qid2')['qid2'].transform('count')
    df['q1len'] = df['question1'].str.len()
    df['q2len'] = df['question2'].str.len()
    df['q1_n_words'] = df['question1'].apply(lambda row: len(row.split(" ")))
    df['q2_n_words'] = df['question2'].apply(lambda row: len(row.split(" ")))

    def normalized_word_Common(row):
        w1 = set(map(lambda word: word.lower().strip(), row['question1'].split(" ")))
        w2 = set(map(lambda word: word.lower().strip(), row['question2'].split(" ")))
        return 1.0 * len(w1 & w2)
    df['word_Common'] = df.apply(normalized_word_Common, axis=1)

    def normalized_word_Total(row):
        w1 = set(map(lambda word: word.lower().strip(), row['question1'].split(" ")))
        w2 = set(map(lambda word: word.lower().strip(), row['question2'].split(" ")))
        return 1.0 * (len(w1) + len(w2))
    df['word_Total'] = df.apply(normalized_word_Total, axis=1)

    def normalized_word_share(row):
        w1 = set(map(lambda word: word.lower().strip(), row['question1'].split(" ")))
        w2 = set(map(lambda word: word.lower().strip(), row['question2'].split(" ")))
        return 1.0 * len(w1 & w2)/(len(w1) + len(w2))
    df['word_share'] = df.apply(normalized_word_share, axis=1)

    df['freq_q1+q2'] = df['freq_qid1']+df['freq_qid2']
    df['freq_q1-q2'] = abs(df['freq_qid1']-df['freq_qid2'])

    df.to_csv("df_fe_without_preprocessing_train.csv", index=False)

df.head()
```

Out[20]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_Total	word_share	freq_q1+q2	freq_q1-q2
0	0	1	2	What is the step by step guide to invest in sh...	What is the step by step guide to invest in sh...	0	1	1	66	57	14	12	10.0	23.0	0.434783	2	0
1	1	3	4	What is the story of Kohinoor (Koh-i-Noor) Dia...	What would happen if the Indian government sto...	0	4	1	51	88	8	13	4.0	20.0	0.200000	5	3
2	2	5	6	How can I increase the speed of my internet co...	How can Internet speed be increased by hacking...	0	1	1	73	59	14	10	4.0	24.0	0.166667	2	0
3	3	7	8	Why am I mentally very lonely? How can I solve...	Find the remainder when [math]23^{24}[/math] i...	0	1	1	50	65	11	9	0.0	19.0	0.000000	2	0
4	4	9	10	Which one dissolve in water quikly sugar, salt...	Which fish would survive in salt water?	0	3	1	76	39	13	7	2.0	20.0	0.100000	4	2

## Analysis of some of the extracted features

- Here are some questions have only one single words.

In [0]:

```
print ("Minimum length of the questions in question1 : " , min(df['q1_n_words']))

print ("Minimum length of the questions in question2 : " , min(df['q2_n_words']))

print ("Number of Questions with minimum length [question1] :", df[df['q1_n_words']== 1].shape[0])
print ("Number of Questions with minimum length [question2] :", df[df['q2_n_words']== 1].shape[0])
```


Minimum length of the questions in question1 : 1  
Minimum length of the questions in question2 : 1  
Number of Questions with minimum length [question1] : 67  
Number of Questions with minimum length [question2] : 24

Feature: word\_share

```
In [0]: plt.figure(figsize=(12, 8))

plt.subplot(1,2,1)
sns.violinplot(x = 'is_duplicate', y = 'word_share', data = df[0:])

plt.subplot(1,2,2)
sns.distplot(df[df['is_duplicate'] == 1.0]['word_share'][0:], label = "1", color = 'red')
sns.distplot(df[df['is_duplicate'] == 0.0]['word_share'][0:], label = "0" , color = 'blue' )
plt.show()
```




- The distributions for normalized word\_share have some overlap on the far right-hand side, i.e., there are quite a lot of questions with high word similarity
- The average word share and Common no. of words of qid1 and qid2 is more when they are duplicate(Similar)

Feature: word\_Common

```
In [0]: plt.figure(figsize=(12, 8))

plt.subplot(1,2,1)
sns.violinplot(x = 'is_duplicate', y = 'word_Common', data = df[0:])

plt.subplot(1,2,2)
sns.distplot(df[df['is_duplicate'] == 1.0]['word_Common'][0:], label = "1", color = 'red')
sns.distplot(df[df['is_duplicate'] == 0.0]['word_Common'][0:], label = "0" , color = 'blue' )
plt.show()
```



The distributions of the word\_Common feature in similar and non-similar questions are highly overlapping

Advanced Feature Extraction.



```
In [0]: import warnings
warnings.filterwarnings("ignore")
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from subprocess import check_output
%matplotlib inline
import plotly.offline as py
py.init_notebook_mode(connected=True)
import plotly.graph_objs as go
import plotly.tools as tls
import os
import gc

import re
from nltk.corpus import stopwords
import distance
from nltk.stem import PorterStemmer
from bs4 import BeautifulSoup
import re
from nltk.corpus import stopwords
# This package is used for finding Longest common subsequence between two strings
# you can write your own dp code for this
import distance
from nltk.stem import PorterStemmer
from bs4 import BeautifulSoup
from fuzzywuzzy import fuzz
from sklearn.manifold import TSNE
# Import the Required Lib packages for WORD-Cloud generation
# https://stackoverflow.com/questions/45625434/how-to-install-wordcloud-in-python3-6
from wordcloud import WordCloud, STOPWORDS
from os import path
from PIL import Image
```

```
In [0]: #https://stackoverflow.com/questions/12468179/unicodedecodeerror-utf8-codec-cant-decode-byte-0x9c
if os.path.isfile('df_fe_without_preprocessing_train.csv'):
    df = pd.read_csv("df_fe_without_preprocessing_train.csv",encoding='latin-1')
    df = df.fillna('')
    df.head()
```

```
In [0]: df.head(2)
```

Out[8]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_Total	word_share	freq_q1+q2	freq_q1-q2
0	0	1	2	What is the step by step guide to invest in sh...	What is the step by step guide to invest in sh...	0	1	1	66	57	14	12	10.0	23.0	0.434783	2	0
1	1	3	4	What is the story of Kohinoor (Koh-i-Noor) Dia...	What would happen if the Indian government sto...	0	4	1	51	88	8	13	4.0	20.0	0.200000	5	3

Preprocessing of Text

- Preprocessing:
  - Removing html tags
  - Removing Punctuations
  - Performing stemming
  - Removing Stopwords
  - Expanding contractions etc.

```
In [0]: # To get the results in 4 decemal points
SAFE_DIV = 0.0001

STOP_WORDS = stopwords.words("english")

def preprocess(x):
    x = str(x).lower()
    x = x.replace(",000,000", "m").replace(",000", "k").replace("'", "").replace("'", "")\
        .replace("won't", "will not").replace("cannot", "can not").replace("can't", "can not")\
        .replace("n't", " not").replace("what's", "what is").replace("it's", "it is")\
        .replace("'ve", " have").replace("i'm", "i am").replace("'re", " are")\
        .replace("he's", "he is").replace("she's", "she is").replace("'s", " own")\
        .replace("%", " percent ").replace("₹", " rupee ").replace("$", " dollar ")\
        .replace("€", " euro ").replace("'ll", " will")

    x = re.sub(r"([0-9]+)000000", r"\1m", x)
    x = re.sub(r"([0-9]+)000", r"\1k", x)

    porter = PorterStemmer()
    pattern = re.compile('\W')

    if type(x) == type(''):
        x = re.sub(pattern, ' ', x)

    if type(x) == type(''):
        x = porter.stem(x)
        example1 = BeautifulSoup(x)
        x = example1.get_text()

    return x
```

- Function to Compute and get the features : With 2 parameters of Question 1 and Question 2

Advanced Feature Extraction (NLP and Fuzzy Features)

Definition:

- **Token**: You get a token by splitting sentence a space
- **Stop\_Word** : stop words as per NLTK.
- **Word** : A token that is not a stop\_word

Features:

- **cwc\_min** : Ratio of common\_word\_count to min length of word count of Q1 and Q2  
cwc\_min = common\_word\_count / (min(len(q1\_words), len(q2\_words)))
- **cwc\_max** : Ratio of common\_word\_count to max length of word count of Q1 and Q2  
cwc\_max = common\_word\_count / (max(len(q1\_words), len(q2\_words)))
- **csc\_min** : Ratio of common\_stop\_count to min length of stop count of Q1 and Q2  
csc\_min = common\_stop\_count / (min(len(q1\_stops), len(q2\_stops)))

- **csc\_max** : Ratio of common\_stop\_count to max lengthth of stop count of Q1 and Q2  
 $csc\_max = common\_stop\_count / (max(len(q1\_stops), len(q2\_stops)))$
- **ctc\_min** : Ratio of common\_token\_count to min lengthth of token count of Q1 and Q2  
 $ctc\_min = common\_token\_count / (min(len(q1\_tokens), len(q2\_tokens)))$
- **ctc\_max** : Ratio of common\_token\_count to max lengthth of token count of Q1 and Q2  
 $ctc\_max = common\_token\_count / (max(len(q1\_tokens), len(q2\_tokens)))$
- **last\_word\_eq** : Check if First word of both questions is equal or not  
 $last\_word\_eq = int(q1\_tokens[-1] == q2\_tokens[-1])$
- **first\_word\_eq** : Check if First word of both questions is equal or not  
 $first\_word\_eq = int(q1\_tokens[0] == q2\_tokens[0])$
- **abs\_len\_diff** : Abs. length difference  
 $abs\_len\_diff = abs(len(q1\_tokens) - len(q2\_tokens))$
- **mean\_len** : Average Token Length of both Questions  
 $mean\_len = (len(q1\_tokens) + len(q2\_tokens))/2$
- **fuzz\_ratio** : <https://github.com/seatgeek/fuzzywuzzy#usage> (<https://github.com/seatgeek/fuzzywuzzy#usage>) <http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>  
(<http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>)
- **fuzz\_partial\_ratio** : <https://github.com/seatgeek/fuzzywuzzy#usage> (<https://github.com/seatgeek/fuzzywuzzy#usage>) <http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>  
(<http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>)
- **token\_sort\_ratio** : <https://github.com/seatgeek/fuzzywuzzy#usage> (<https://github.com/seatgeek/fuzzywuzzy#usage>) <http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>  
(<http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>)
- **token\_set\_ratio** : <https://github.com/seatgeek/fuzzywuzzy#usage> (<https://github.com/seatgeek/fuzzywuzzy#usage>) <http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>  
(<http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>)
- **longest\_substr\_ratio** : Ratio of length longest common substring to min lengthth of token count of Q1 and Q2  
 $longest\_substr\_ratio = len(longest\ common\ substring) / (min(len(q1\_tokens), len(q2\_tokens)))$

```
In [0]: def get_token_features(q1, q2):
    token_features = [0.0]*10

    # Converting the Sentence into Tokens:
    q1_tokens = q1.split()
    q2_tokens = q2.split()

    if len(q1_tokens) == 0 or len(q2_tokens) == 0:
        return token_features
    # Get the non-stopwords in Questions
    q1_words = set([word for word in q1_tokens if word not in STOP_WORDS])
    q2_words = set([word for word in q2_tokens if word not in STOP_WORDS])

    #Get the stopwords in Questions
    q1_stops = set([word for word in q1_tokens if word in STOP_WORDS])
    q2_stops = set([word for word in q2_tokens if word in STOP_WORDS])

    # Get the common non-stopwords from Question pair
    common_word_count = len(q1_words.intersection(q2_words))

    # Get the common stopwords from Question pair
    common_stop_count = len(q1_stops.intersection(q2_stops))

    # Get the common Tokens from Question pair
    common_token_count = len(set(q1_tokens).intersection(set(q2_tokens)))

    token_features[0] = common_word_count / (min(len(q1_words), len(q2_words)) + SAFE_DIV)
    token_features[1] = common_word_count / (max(len(q1_words), len(q2_words)) + SAFE_DIV)
    token_features[2] = common_stop_count / (min(len(q1_stops), len(q2_stops)) + SAFE_DIV)
    token_features[3] = common_stop_count / (max(len(q1_stops), len(q2_stops)) + SAFE_DIV)
    token_features[4] = common_token_count / (min(len(q1_tokens), len(q2_tokens)) + SAFE_DIV)
    token_features[5] = common_token_count / (max(len(q1_tokens), len(q2_tokens)) + SAFE_DIV)

    # Last word of both question is same or not
    token_features[6] = int(q1_tokens[-1] == q2_tokens[-1])

    # First word of both question is same or not
    token_features[7] = int(q1_tokens[0] == q2_tokens[0])

    token_features[8] = abs(len(q1_tokens) - len(q2_tokens))

    #Average Token Length of both Questions
    token_features[9] = (len(q1_tokens) + len(q2_tokens))/2
    return token_features

# get the Longest Common sub string

def get_longest_substr_ratio(a, b):
    strs = list(distance.lcs substrings(a, b))
    if len(strs) == 0:
        return 0
    else:
        return len(strs[0]) / (min(len(a), len(b)) + 1)

def extract_features(df):
    # preprocessing each question
    df["question1"] = df["question1"].fillna("").apply(preprocess)
    df["question2"] = df["question2"].fillna("").apply(preprocess)

    print("token features...")

    # Merging Features with dataset

    token_features = df.apply(lambda x: get_token_features(x["question1"], x["question2"]), axis=1)

    df["cwc_min"] = list(map(lambda x: x[0], token_features))
    df["cwc_max"] = list(map(lambda x: x[1], token_features))
    df["csc_min"] = list(map(lambda x: x[2], token_features))
    df["csc_max"] = list(map(lambda x: x[3], token_features))
    df["ctc_min"] = list(map(lambda x: x[4], token_features))
    df["ctc_max"] = list(map(lambda x: x[5], token_features))
    df["last_word_eq"] = list(map(lambda x: x[6], token_features))
    df["first_word_eq"] = list(map(lambda x: x[7], token_features))
    df["abs_len_diff"] = list(map(lambda x: x[8], token_features))
    df["mean_len"] = list(map(lambda x: x[9], token_features))

    #Computing Fuzzy Features and Merging with Dataset

    # do read this blog: http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/
    # https://stackoverflow.com/questions/31806695/when-to-use-which-fuzz-function-to-compare-2-strings
    # https://github.com/seatgeek/fuzzywuzzy
    print("fuzzy features..")

    df["token_set_ratio"] = df.apply(lambda x: fuzz.token_set_ratio(x["question1"], x["question2"]), axis=1)
    # The token sort approach involves tokenizing the string in question, sorting the tokens alphabetically, and
    # then joining them back into a string We then compare the transformed strings with a simple ratio().
    df["token_sort_ratio"] = df.apply(lambda x: fuzz.token_sort_ratio(x["question1"], x["question2"]), axis=1)
    df["fuzz_ratio"] = df.apply(lambda x: fuzz.QRatio(x["question1"], x["question2"]), axis=1)
    df["fuzz_partial_ratio"] = df.apply(lambda x: fuzz.partial_ratio(x["question1"], x["question2"]), axis=1)
    df["longest_substr_ratio"] = df.apply(lambda x: get_longest_substr_ratio(x["question1"], x["question2"]), axis=1)
    return df
```

```
In [0]: if os.path.isfile('nlp_features_train.csv'):
    df = pd.read_csv("nlp_features_train.csv",encoding='latin-1')
    df.fillna('')
else:
    print("Extracting features for train:")
    df = pd.read_csv("train.csv")
    df = extract_features(df)
    df.to_csv("nlp_features_train.csv", index=False)
df.head(2)
```

Out[12]:

	id	qid1	qid2	question1	question2	is_duplicate	cwc_min	cwc_max	csc_min	csc_max	...	ctc_max	last_word_eq	first_word_eq	abs_len_diff	mean_len	token_set_ratio	token_sort_ratio	fuzz_ratio	fuzz_partial_ratic
0	0	1	2	what is the step by step guide to invest in sh...	what is the step by step guide to invest in sh...	0	0.999980	0.833319	0.999983	0.999983	...	0.785709	0.0	1.0	2.0	13.0	100	93	93	100
1	1	3	4	what is the story of kohinoor koh i noor dia...	what would happen if the indian government sto...	0	0.799984	0.399996	0.749981	0.599988	...	0.466664	0.0	1.0	5.0	12.5	86	63	66	71

2 rows × 21 columns

### Analysis of extracted features

#### Plotting Word clouds



- Creating Word Cloud of Duplicates and Non-Duplicates Question pairs
- We can observe the most frequent occurring words

```
In [0]: df_duplicate = df[df['is_duplicate'] == 1]
dfp_nonduplicate = df[df['is_duplicate'] == 0]

# Converting 2d array of q1 and q2 and flatten the array: Like {{1,2},{3,4}} to {1,2,3,4}
p = np.dstack([df_duplicate["question1"], df_duplicate["question2"]]).flatten()
n = np.dstack([dfp_nonduplicate["question1"], dfp_nonduplicate["question2"]]).flatten()

print ("Number of data points in class 1 (duplicate pairs) :",len(p))
print ("Number of data points in class 0 (non duplicate pairs) :",len(n))

#Saving the np array into a text file
np.savetxt('train_p.txt', p, delimiter=' ', fmt='%s')
np.savetxt('train_n.txt', n, delimiter=' ', fmt='%s')
```

```
Number of data points in class 1 (duplicate pairs) : 298526
Number of data points in class 0 (non duplicate pairs) : 510054
```

```
In [0]: # reading the text files and removing the Stop Words:
d = path.dirname('.')

textp_w = open(path.join(d, 'train_p.txt')).read()
textn_w = open(path.join(d, 'train_n.txt')).read()
stopwords = set(STOPWORDS)
stopwords.add("said")
stopwords.add("br")
stopwords.add(" ")
stopwords.remove("not")

stopwords.remove("no")
#stopwords.remove("good")
#stopwords.remove("Love")
stopwords.remove("like")
#stopwords.remove("best")
#stopwords.remove("!")
print("Total number of words in duplicate pair questions :",len(textp_w))
print("Total number of words in non duplicate pair questions :",len(textn_w))
```

Total number of words in duplicate pair questions : 16109886  
Total number of words in non duplicate pair questions : 33193130

**Word Clouds generated from duplicate pair question's text**

```
In [0]: wc = WordCloud(background_color="white", max_words=len(textp_w), stopwords=stopwords)
wc.generate(textp_w)
print ("Word Cloud for Duplicate Question pairs")
plt.imshow(wc, interpolation='bilinear')
plt.axis("off")
plt.show()
```


Word Cloud for Duplicate Question pairs



**Word Clouds generated from non duplicate pair question's text**

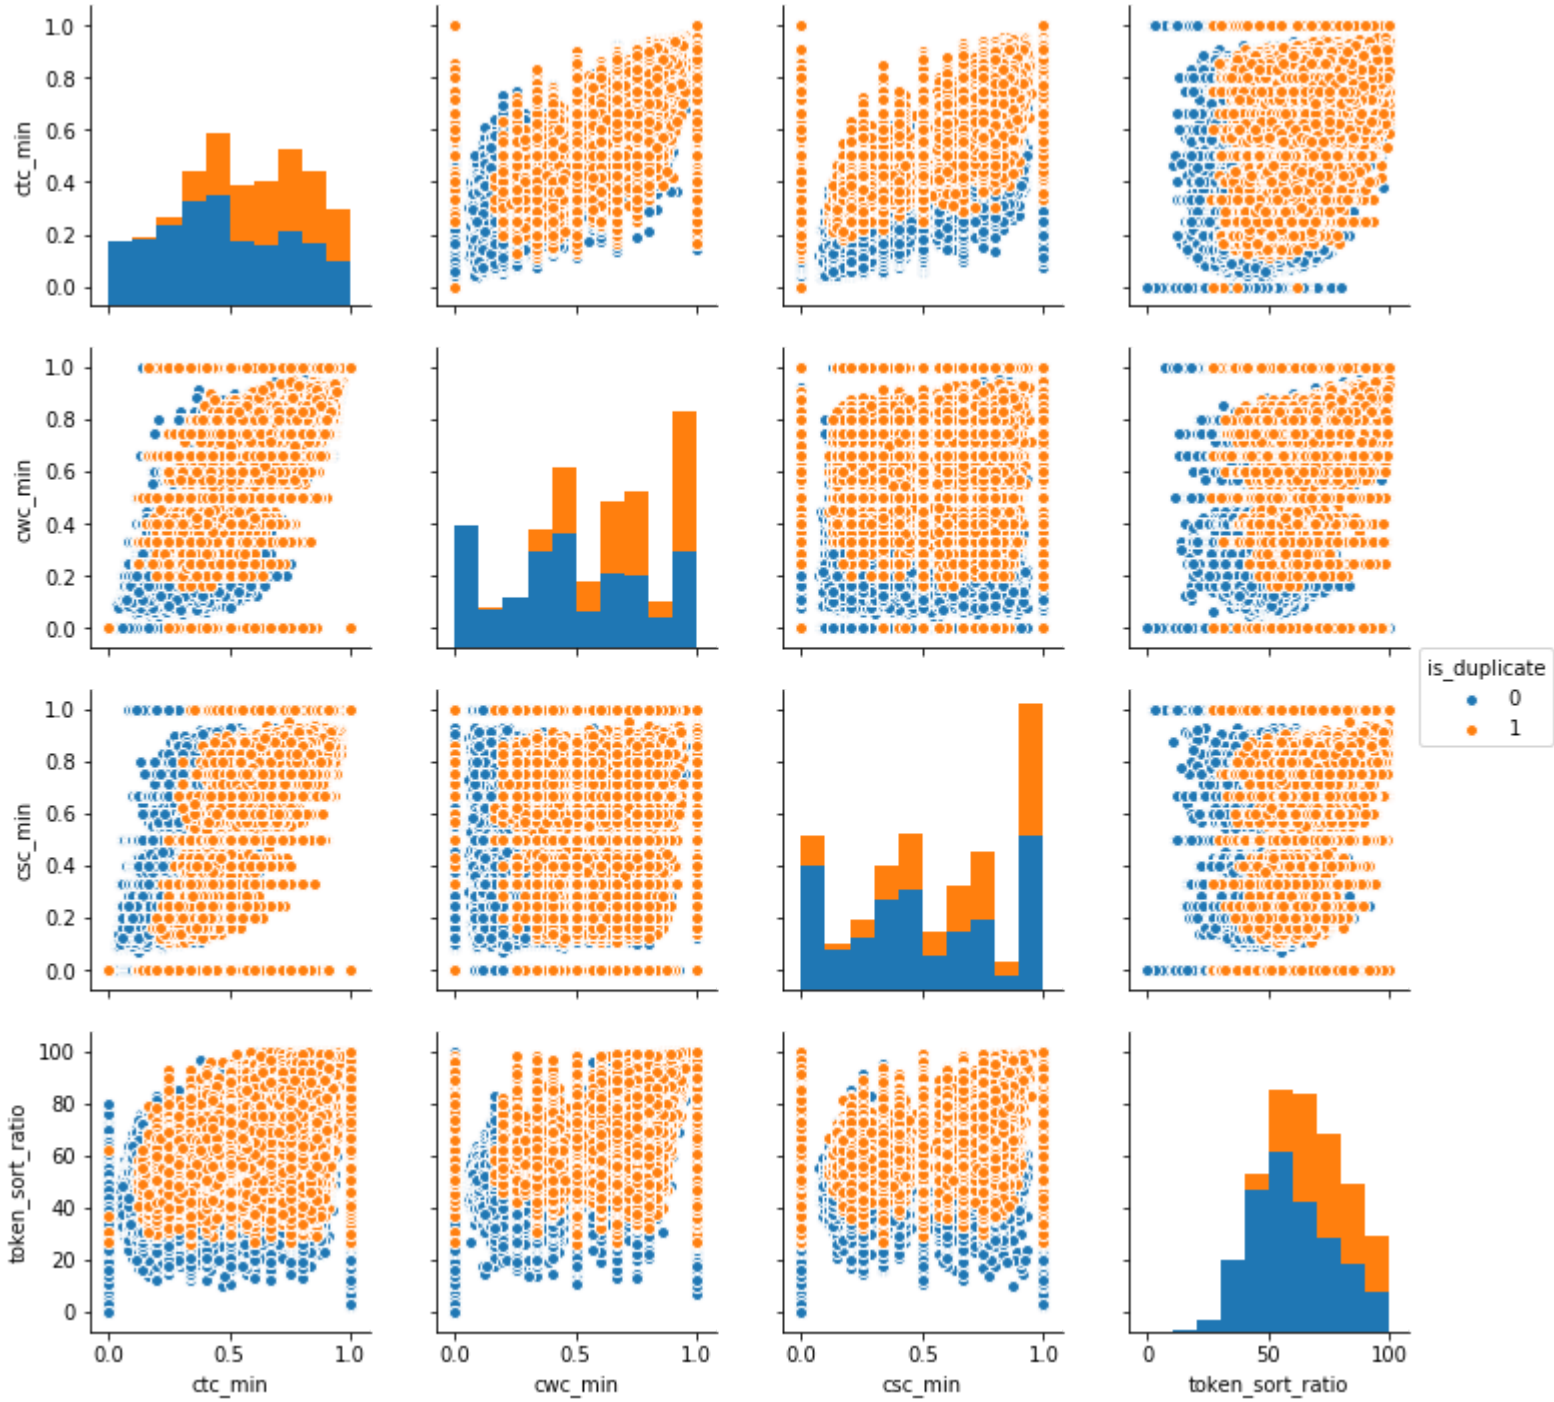
```
In [0]: wc = WordCloud(background_color="white", max_words=len(textn_w), stopwords=stopwords)
# generate word cloud
wc.generate(textn_w)
print("Word Cloud for non-Duplicate Question pairs:")
plt.imshow(wc, interpolation='bilinear')
plt.axis("off")
plt.show()
```

Word Cloud for non-Duplicate Question pairs:



### Pair plot of features ['ctc\_min', 'cwc\_min', 'csc\_min', 'token\_sort\_ratio']

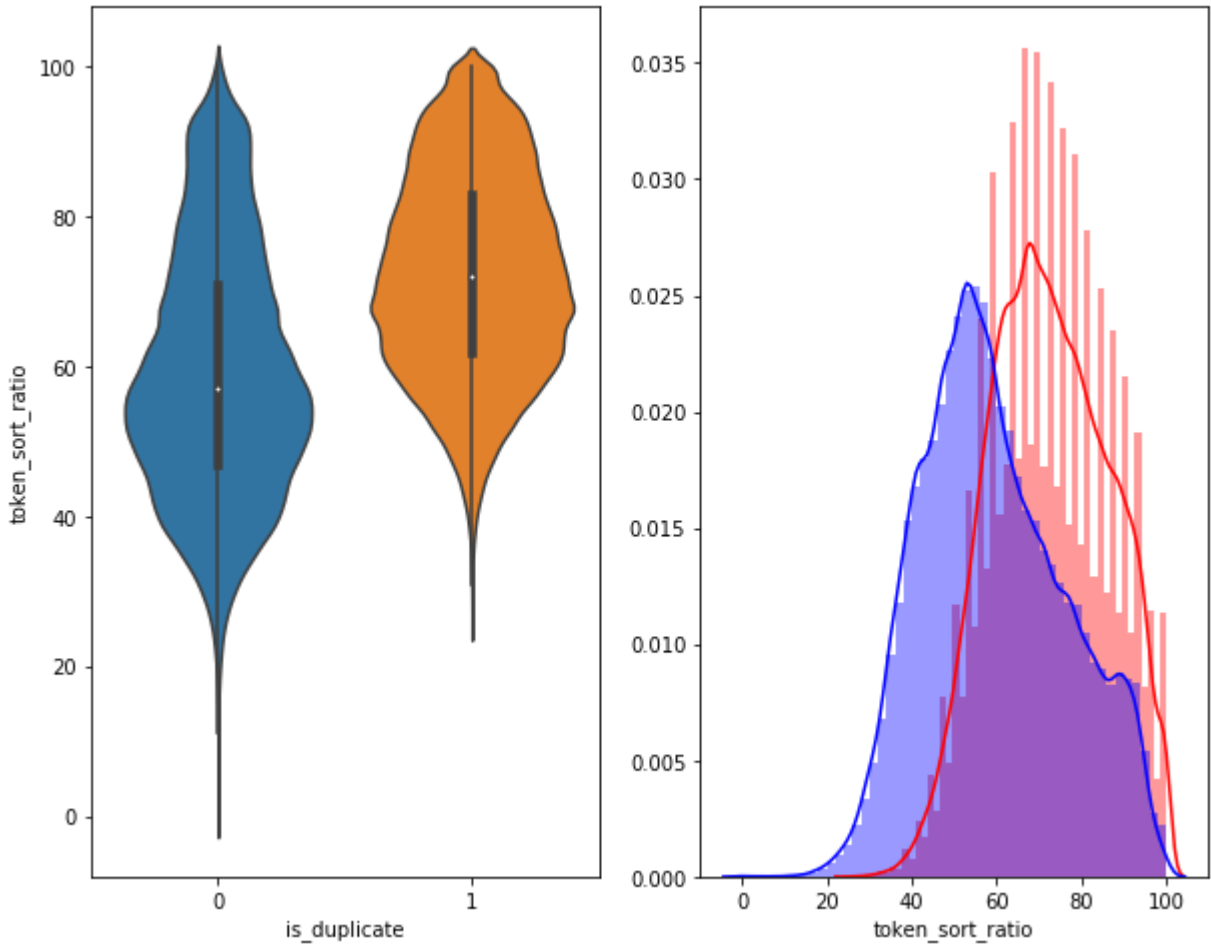
```
In [0]: n = df.shape[0]
sns.pairplot(df[['ctc_min', 'cwc_min', 'csc_min', 'token_sort_ratio', 'is_duplicate']][0:n], hue='is_duplicate', vars=['ctc_min', 'cwc_min', 'csc_min', 'token_sort_ratio'])
plt.show()
```



```
In [0]: # Distribution of the token_sort_ratio
plt.figure(figsize=(10, 8))

plt.subplot(1,2,1)
sns.violinplot(x = 'is_duplicate', y = 'token_sort_ratio', data = df[0:] , )

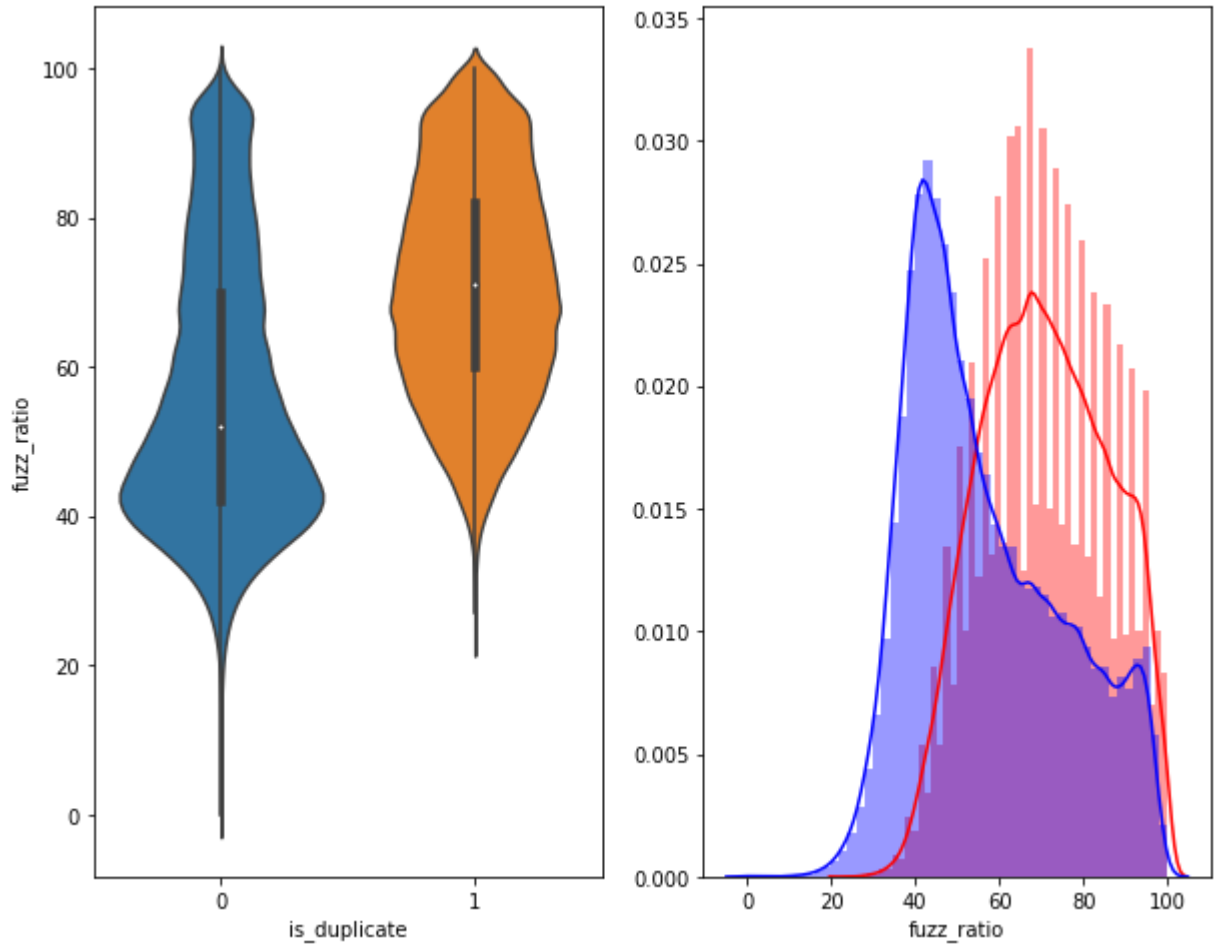
plt.subplot(1,2,2)
sns.distplot(df[df['is_duplicate'] == 1.0]['token_sort_ratio'][0:] , label = "1", color = 'red')
sns.distplot(df[df['is_duplicate'] == 0.0]['token_sort_ratio'][0:] , label = "0" , color = 'blue' )
plt.show()
```



```
In [0]: plt.figure(figsize=(10, 8))

plt.subplot(1,2,1)
sns.violinplot(x = 'is_duplicate', y = 'fuzz_ratio', data = df[0:] , )

plt.subplot(1,2,2)
sns.distplot(df[df['is_duplicate'] == 1.0]['fuzz_ratio'][0:] , label = "1", color = 'red')
sns.distplot(df[df['is_duplicate'] == 0.0]['fuzz_ratio'][0:] , label = "0" , color = 'blue' )
plt.show()
```



Visualization

```
In [0]: # Using TSNE for Dimentionality reduction for 15 Features(Generated after cLeaning the data) to 3 dimention

from sklearn.preprocessing import MinMaxScaler

dfp_subsampled = df[0:5000]
X = MinMaxScaler().fit_transform(dfp_subsampled[['cwc_min', 'cwc_max', 'csc_min', 'csc_max' , 'ctc_min' , 'ctc_max' , 'last_word_eq', 'first_word_eq' , 'abs_len_diff' , 'mean_len' ,
y = dfp_subsampled['is_duplicate'].values
```


```
In [0]: tsne2d = TSNE(
    n_components=2,
    init='random', # pca
    random_state=101,
    method='barnes_hut',
    n_iter=1000,
    verbose=2,
    angle=0.5
).fit_transform(X)

[t-SNE] Computing 91 nearest neighbors...
[t-SNE] Indexed 5000 samples in 0.011s...
[t-SNE] Computed neighbors for 5000 samples in 0.912s...
[t-SNE] Computed conditional probabilities for sample 1000 / 5000
[t-SNE] Computed conditional probabilities for sample 2000 / 5000
[t-SNE] Computed conditional probabilities for sample 3000 / 5000
[t-SNE] Computed conditional probabilities for sample 4000 / 5000
[t-SNE] Computed conditional probabilities for sample 5000 / 5000
[t-SNE] Mean sigma: 0.116557
[t-SNE] Computed conditional probabilities in 0.433s
[t-SNE] Iteration 50: error = 80.9244080, gradient norm = 0.0428133 (50 iterations in 13.099s)
[t-SNE] Iteration 100: error = 70.3858795, gradient norm = 0.0100968 (50 iterations in 9.067s)
[t-SNE] Iteration 150: error = 68.6138382, gradient norm = 0.0058392 (50 iterations in 9.602s)
[t-SNE] Iteration 200: error = 67.7700119, gradient norm = 0.0036596 (50 iterations in 9.121s)
[t-SNE] Iteration 250: error = 67.2725067, gradient norm = 0.0034962 (50 iterations in 11.305s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 67.272507
[t-SNE] Iteration 300: error = 1.7737305, gradient norm = 0.0011918 (50 iterations in 8.289s)
[t-SNE] Iteration 350: error = 1.3720417, gradient norm = 0.0004822 (50 iterations in 10.526s)
[t-SNE] Iteration 400: error = 1.2039998, gradient norm = 0.0002768 (50 iterations in 9.600s)
[t-SNE] Iteration 450: error = 1.1133438, gradient norm = 0.0001881 (50 iterations in 11.827s)
[t-SNE] Iteration 500: error = 1.0579143, gradient norm = 0.0001434 (50 iterations in 8.941s)
[t-SNE] Iteration 550: error = 1.0221983, gradient norm = 0.0001164 (50 iterations in 11.092s)
[t-SNE] Iteration 600: error = 0.9987167, gradient norm = 0.0001039 (50 iterations in 11.467s)
[t-SNE] Iteration 650: error = 0.9831534, gradient norm = 0.0000938 (50 iterations in 11.799s)
[t-SNE] Iteration 700: error = 0.9722011, gradient norm = 0.0000858 (50 iterations in 12.028s)
[t-SNE] Iteration 750: error = 0.9643636, gradient norm = 0.0000799 (50 iterations in 12.120s)
[t-SNE] Iteration 800: error = 0.9584482, gradient norm = 0.0000785 (50 iterations in 11.867s)
[t-SNE] Iteration 850: error = 0.9538348, gradient norm = 0.0000739 (50 iterations in 11.461s)
[t-SNE] Iteration 900: error = 0.9496906, gradient norm = 0.0000712 (50 iterations in 11.023s)
[t-SNE] Iteration 950: error = 0.9463405, gradient norm = 0.0000673 (50 iterations in 11.755s)
[t-SNE] Iteration 1000: error = 0.9432716, gradient norm = 0.0000662 (50 iterations in 11.493s)
[t-SNE] Error after 1000 iterations: 0.943272
```



```
In [0]: df = pd.DataFrame({'x':tsne2d[:,0], 'y':tsne2d[:,1] , 'label':y})

# draw the plot in appropriate place in the grid
sns.lmplot(data=df, x='x', y='y', hue='label', fit_reg=False, size=8,palette="Set1",markers=['s','o'])
plt.title("perplexity : {} and max_iter : {}".format(30, 1000))
plt.show()
```



```
In [0]: from sklearn.manifold import TSNE
tsne3d = TSNE(
    n_components=3,
    init='random', # pca
    random_state=101,
    method='barnes_hut',
    n_iter=1000,
    verbose=2,
    angle=0.5
).fit_transform(X)

[t-SNE] Computing 91 nearest neighbors...
[t-SNE] Indexed 5000 samples in 0.010s...
[t-SNE] Computed neighbors for 5000 samples in 0.935s...
[t-SNE] Computed conditional probabilities for sample 1000 / 5000
[t-SNE] Computed conditional probabilities for sample 2000 / 5000
[t-SNE] Computed conditional probabilities for sample 3000 / 5000
[t-SNE] Computed conditional probabilities for sample 4000 / 5000
[t-SNE] Computed conditional probabilities for sample 5000 / 5000
[t-SNE] Mean sigma: 0.116557
[t-SNE] Computed conditional probabilities in 0.363s
[t-SNE] Iteration 50: error = 77.7944183, gradient norm = 0.1014017 (50 iterations in 34.931s)
[t-SNE] Iteration 100: error = 69.2682266, gradient norm = 0.0248657 (50 iterations in 15.147s)
[t-SNE] Iteration 150: error = 67.7877655, gradient norm = 0.0150941 (50 iterations in 13.761s)
[t-SNE] Iteration 200: error = 67.1991119, gradient norm = 0.0126559 (50 iterations in 13.425s)
[t-SNE] Iteration 250: error = 66.8560715, gradient norm = 0.0074975 (50 iterations in 12.904s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 66.856071
[t-SNE] Iteration 300: error = 1.2356015, gradient norm = 0.0007033 (50 iterations in 13.302s)
[t-SNE] Iteration 350: error = 0.9948602, gradient norm = 0.0001997 (50 iterations in 18.898s)
[t-SNE] Iteration 400: error = 0.9168936, gradient norm = 0.0001430 (50 iterations in 13.397s)
[t-SNE] Iteration 450: error = 0.8863022, gradient norm = 0.0000975 (50 iterations in 16.379s)
[t-SNE] Iteration 500: error = 0.8681002, gradient norm = 0.0000854 (50 iterations in 17.791s)
[t-SNE] Iteration 550: error = 0.8564141, gradient norm = 0.0000694 (50 iterations in 17.060s)
[t-SNE] Iteration 600: error = 0.8470711, gradient norm = 0.0000640 (50 iterations in 15.454s)
[t-SNE] Iteration 650: error = 0.8389117, gradient norm = 0.0000561 (50 iterations in 17.562s)
[t-SNE] Iteration 700: error = 0.8325295, gradient norm = 0.0000529 (50 iterations in 13.443s)
[t-SNE] Iteration 750: error = 0.8268463, gradient norm = 0.0000528 (50 iterations in 17.981s)
[t-SNE] Iteration 800: error = 0.8219477, gradient norm = 0.0000477 (50 iterations in 17.448s)
[t-SNE] Iteration 850: error = 0.8180174, gradient norm = 0.0000490 (50 iterations in 18.376s)
[t-SNE] Iteration 900: error = 0.8150476, gradient norm = 0.0000456 (50 iterations in 17.778s)
[t-SNE] Iteration 950: error = 0.8122067, gradient norm = 0.0000472 (50 iterations in 16.983s)
[t-SNE] Iteration 1000: error = 0.8095787, gradient norm = 0.0000489 (50 iterations in 18.581s)
[t-SNE] Error after 1000 iterations: 0.809579
```

```
In [0]: trace1 = go.Scatter3d(
    x=tsne3d[:,0],
    y=tsne3d[:,1],
    z=tsne3d[:,2],
    mode='markers',
    marker=dict(
        sizemode='diameter',
        color = y,
        colorscale = 'Portland',
        colorbar = dict(title = 'duplicate'),
        line=dict(color='rgb(255, 255, 255)'),
        opacity=0.75
    )
)

data=[trace1]
layout=dict(height=800, width=800, title='3d embedding with engineered features')
fig=dict(data=data, layout=layout)
py.iplot(fig, filename='3DBubble')
```

## Text Featurization I

### TFIDF

```
In [0]: #prepro_features_train.csv (Simple Preprocessing Feartures)
#nlp_features_train.csv (NLP Features)
if os.path.isfile('Quora_question_pair_similarity/data/nlp_features_train.csv'):
    dfnlp = pd.read_csv("Quora_question_pair_similarity/data/nlp_features_train.csv",encoding='latin-1')

if os.path.isfile('Quora_question_pair_similarity/data/df_fe_without_preprocessing_train.csv'):
    dfppro = pd.read_csv("Quora_question_pair_similarity/data/df_fe_without_preprocessing_train.csv",encoding='latin-1')
```

```
In [0]: dfppro.head()
```

Out[7]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_Total	word_share	freq_q1+q2	freq_q1-q2
0	0	1	2	What is the step by step guide to invest in sh...	What is the step by step guide to invest in sh...	0	1	1	66	57	14	12	10.0	23.0	0.434783	2	0
1	1	3	4	What is the story of Kohinoor (Koh-i-Noor) Dia...	What would happen if the Indian government sto...	0	4	1	51	88	8	13	4.0	20.0	0.200000	5	3
2	2	5	6	How can I increase the speed of my internet co...	How can Internet speed be increased by hacking...	0	1	1	73	59	14	10	4.0	24.0	0.166667	2	0
3	3	7	8	Why am I mentally very lonely? How can I solve...	Find the remainder when $23^{24}$ is divided by 29	0	1	1	50	65	11	9	0.0	19.0	0.000000	2	0
4	4	9	10	Which one dissolve in water quikly sugar, salt...	Which fish would survive in salt water?	0	3	1	76	39	13	7	2.0	20.0	0.100000	4	2



In [0]: dfnlp.head()

Out[8]:

	id	qid1	qid2	question1	question2	is_duplicate	cwc_min	cwc_max	csc_min	csc_max	...	ctc_max	last_word_eq	first_word_eq	abs_len_diff	mean_len	token_set_ratio	token_sort_ratio	fuzz_ratio	fuzz_partial_rati
0	0	1	2	what is the step by step guide to invest in sh...	what is the step by step guide to invest in sh...	0	0.999980	0.833319	0.999983	0.999983	...	0.785709	0.0	1.0	2.0	13.0	100	93	93	100
1	1	3	4	what is the story of kohinoor koh i noor dia...	what would happen if the indian government sto...	0	0.799984	0.399996	0.749981	0.599988	...	0.466664	0.0	1.0	5.0	12.5	86	63	66	71
2	2	5	6	how can i increase the speed of my internet co...	how can internet speed be increased by hacking...	0	0.399992	0.333328	0.399992	0.249997	...	0.285712	0.0	1.0	4.0	12.0	66	66	54	54
3	3	7	8	why am i mentally very lonely how can i solve...	find the remainder when math 23 24 math i...	0	0.000000	0.000000	0.000000	0.000000	...	0.000000	0.0	0.0	2.0	12.0	36	36	35	40
4	4	9	10	which one dissolve in water quickly sugar salt...	which fish would survive in salt water	0	0.399992	0.199998	0.999950	0.666644	...	0.307690	0.0	1.0	6.0	10.0	67	47	46	50

5 rows × 21 columns

In [0]: print('The shape of dfppro is',dfppro.shape)  
print('The shape of dfnlp is',dfnlp.shape)

The shape of dfppro is (404290, 17)  
The shape of dfnlp is (404290, 21)

In [0]: df\_feats\_nlp = dfnlp.drop(['qid1','qid2'],axis=1)  
df\_feats\_pro = dfppro.drop(['qid1','qid2','question1','question2','is\_duplicate'],axis=1)

In [0]: df = df\_feats\_nlp.merge(df\_feats\_pro, on='id',how='left')  
print('Shape of final merged data',df.shape)

Shape of final merged data (404290, 30)

In [0]: df.head()

Out[14]:

	id	question1	question2	is_duplicate	cwc_min	cwc_max	csc_min	csc_max	ctc_min	ctc_max	...	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_Total	word_share	freq_q1+q2	freq_q1-q2
0	0	what is the step by step guide to invest in sh...	what is the step by step guide to invest in sh...	0	0.999980	0.833319	0.999983	0.999983	0.916659	0.785709	...	1	66	57	14	12	10.0	23.0	0.434783	2	0
1	1	what is the story of kohinoor koh i noor dia...	what would happen if the indian government sto...	0	0.799984	0.399996	0.749981	0.599988	0.699993	0.466664	...	1	51	88	8	13	4.0	20.0	0.200000	5	3
2	2	how can i increase the speed of my internet co...	how can internet speed be increased by hacking...	0	0.399992	0.333328	0.399992	0.249997	0.399996	0.285712	...	1	73	59	14	10	4.0	24.0	0.166667	2	0
3	3	why am i mentally very lonely how can i solve...	find the remainder when math 23 24 math i...	0	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	...	1	50	65	11	9	0.0	19.0	0.000000	2	0
4	4	which one dissolve in water quickly sugar salt...	which fish would survive in salt water	0	0.399992	0.199998	0.999950	0.666644	0.571420	0.307690	...	1	76	39	13	7	2.0	20.0	0.100000	4	2

5 rows × 30 columns

In [0]: df['is\_duplicate'].value\_counts()

Out[13]:

```
0    255027
1    149263
Name: is_duplicate, dtype: int64
```

Observations: We can see that 255027 question pairs that are duplicate, 149263 question pairs that are not duplicate.

In [0]: y\_feats = df['is\_duplicate']  
df.drop(['id','is\_duplicate'],axis=1,inplace=True)  
  
x\_train,x\_test,y\_train,y\_test = train\_test\_split(df, y\_feats, stratify=y\_feats, test\_size=0.3)  
  
print('The shape of x\_train is {}'.format(x\_train.shape))  
print('The shape of y\_train is {}'.format(y\_train.shape))  
print('The shape of x\_test is {}'.format(x\_test.shape))  
print('The shape of y\_test is {}'.format(y\_test.shape))

The shape of x\_train is (283003, 28)  
The shape of y\_train is (283003,)  
The shape of x\_test is (121287, 28)  
The shape of y\_test is (121287,)

In [0]: y\_train.value\_counts()

Out[15]:

```
0    178519
1    104484
Name: is_duplicate, dtype: int64
```

```
In [0]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer

tfidf_q1 = TfidfVectorizer() #TFIDF vectorizer for question 1
train_q1 = tfidf_q1.fit_transform(x_train['question1'].values.astype('U'))
test_q1 = tfidf_q1.transform(x_test['question1'].values.astype('U'))

tfidf_q2 = TfidfVectorizer() #TFIDF Vectorizer for question 2
train_q2 = tfidf_q2.fit_transform(x_train['question2'].values.astype('U'))
test_q2 = tfidf_q2.transform(x_test['question2'].values.astype('U'))

In [0]: train_feats = hstack((train_q1, train_q2)) #Stack train tfidf vectors for question 1 and question 2
test_feats = hstack((test_q1, test_q2)) #Stack test tfidf vectors for question 1 and question 2

print('The shape of TFIDF vectorization of train data is',train_feats.shape)
print('The shape of TFIDF vectorization of test data is',test_feats.shape)

The shape of TFIDF vectorization of train data is (283003, 111664)
The shape of TFIDF vectorization of test data is (121287, 111664)

In [0]: x_train.drop(['question1','question2'], axis=1, inplace=True) #Drop unwanted columns
x_test.drop(['question1','question2'], axis=1, inplace=True)

In [0]: X_train=x_train.as_matrix() #Convert to matrix
X_test=x_test.as_matrix()

In [0]: X_train = hstack((X_train,train_feats)) #Finally feature engineering with tfidf features of question 1 and 2
X_test = hstack((X_test,test_feats))
```

# Machine Learning Models I

## Random Model

```
In [0]: predicted_y = np.zeros((len(y_test),2))
for i in range(len(y_test)):
    rand_probs = np.random.rand(1,2)
    predicted_y[i] = ((rand_probs/sum(sum(rand_probs)))[0])
print("Random model Log Loss on Test data is ",log_loss(y_test, predicted_y, eps=1e-15))

predicted_y =np.argmax(predicted_y, axis=1)

Random model Log Loss on Test data is  0.884447493006593
```

## Logistic Regression

```
In [0]: alpha = [0.0001,0.001,0.01,0.1,1,10,100,1000]

predict_train = []
predict_test = []
log_error_test_array=[]
log_error_train_array=[]

for i in alpha:
    clf = SGDClassifier(alpha=i, loss='log')
    clf.fit(X_train, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train, y_train)
    predict_y_test = sig_clf.predict_proba(X_test)
    predict_y_train = sig_clf.predict_proba(X_train)
    predict_test.append(clf.predict(X_test))
    predict_train.append(clf.predict(X_train))
    log_error_test_array.append(log_loss(y_test, predict_y_test, labels=clf.classes_, eps=1e-15))
    log_error_train_array.append(log_loss(y_train, predict_y_train, labels=clf.classes_, eps=1e-15))
    print('When alpha is {} log loss is {}'.format(i,log_loss(y_test, predict_y_test, labels=clf.classes_, eps=1e-15)))

fig, ax = plt.subplots()
ax.plot(alpha, log_error_test_array,c='g')
for i, txt in enumerate(np.round(log_error_test_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_test_array[i]))
plt.grid(linestyle = '-')
plt.title("Cross Validation Error")
plt.xlabel("Alpha")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(log_error_test_array)

print('For values of best alpha {} the train log loss is {}'.format(alpha[best_alpha],log_error_train_array[best_alpha]))

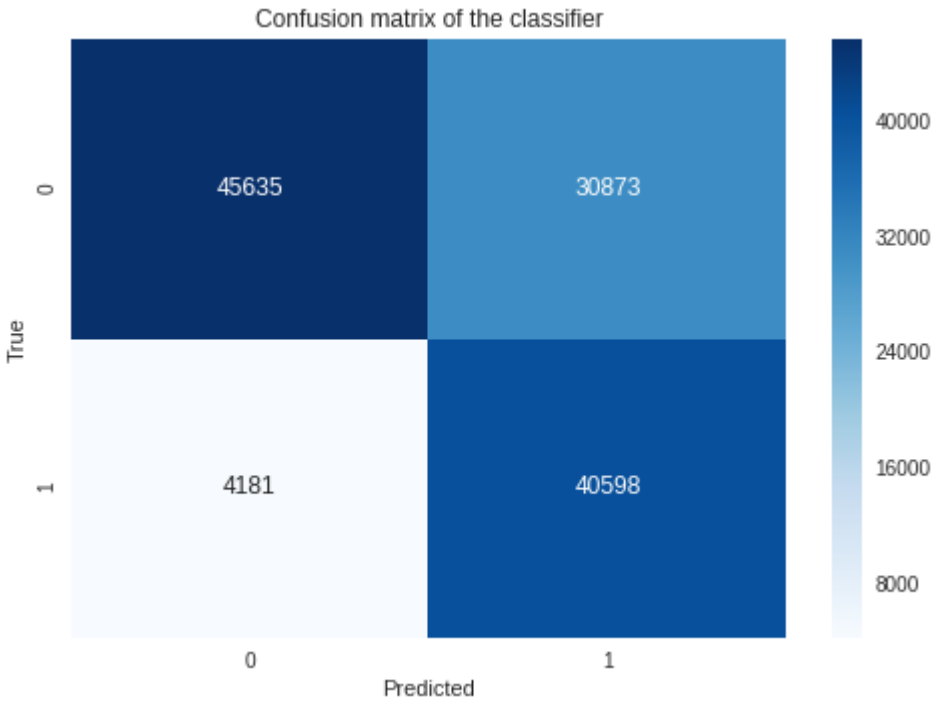
print('For values of best alpha {} the test log loss is {}'.format(alpha[best_alpha],log_error_test_array[best_alpha]))
```

When alpha is 0.0001 log loss is 0.44702936689993683  
When alpha is 0.001 log loss is 0.4592529411350289  
When alpha is 0.01 log loss is 0.44236647075139723  
When alpha is 0.1 log loss is 0.4662027092416626  
When alpha is 1 log loss is 0.4932871094479345  
When alpha is 10 log loss is 0.547060351646075  
When alpha is 100 log loss is 0.5875748231327537  
When alpha is 1000 log loss is 0.6225942575558097



For values of best alpha 0.01 the train log loss is 0.4409423224089968  
For values of best alpha 0.01 the test log loss is 0.44236647075139723

```
In [0]: from pandas_ml import ConfusionMatrix
y_true = np.array(y_test)
y_pred = predict_test[best_alpha]
#print(confusion_matrix(y_test, y_pred))
cm = ConfusionMatrix(y_test,y_pred) #This the confusion matrix of pandas_ml which provides interesting s
confusion_matrix_plot = confusion_matrix(y_test,y_pred) #We are plotting confusion matrix of sklearn
heatmap = sns.heatmap(confusion_matrix_plot, annot=True,cmap='Blues', fmt='g')
plt.title('Confusion matrix of the classifier')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
print(""*50)
print("The True Positive Rate observed is",cm.TPR) #This prints the True Positive Rate of the confusion matrix (using pandas_ml confusion matrix).
print("The True Negative Rate observed is",cm.TNR)
print("The False Positive Rate observed is",cm.FPR)
print("The False Negative Rate observed is",cm.FNR)
print(""*50)
print("The stats observed for confusion matrix are:")
cm.print_stats()#Prints all the stats of the confusion matrix plotted (using pandas_ml confusion matrix)
```



\*\*\*\*\*

The True Positive Rate observed is 0.5926146515783204  
The True Negative Rate observed is 0.4156324267229997  
The False Positive Rate observed is 0.5843675732770003  
The False Negative Rate observed is 0.40738534842167956  
\*\*\*\*\*

The stats observed for confusion matrix are:  
population: 36154  
P: 13432  
N: 22722  
PositiveTest: 21238  
NegativeTest: 14916  
TP: 7960  
TN: 9444  
FP: 13278  
FN: 5472  
TPR: 0.5926146515783204  
TNR: 0.4156324267229997  
PPV: 0.37479988699500894  
NPV: 0.6331456154465004  
FPR: 0.5843675732770003  
FDR: 0.625200113004991  
FNR: 0.40738534842167956  
ACC: 0.4813851855949549  
F1\_score: 0.4591866166714739  
MCC: 0.008094886088047079  
informedness: 0.008247078301319988  
markedness: 0.007945502441509378  
prevalence: 0.37152182331139016  
LRP: 1.0141128267180746  
LRN: 0.9801577601479673  
DOR: 1.0346424503796015  
FOR: 0.3668543845534996

Linear SVM

```
In [0]: alpha = [0.0001,0.001,0.01,0.1,1,10,100,1000]

predict_train = []
predict_test = []
log_error_test_array=[]
log_error_train_array=[]

for i in alpha:
    clf = SGDClassifier(alpha=i, loss='hinge')
    clf.fit(X_train, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train, y_train)
    predict_y_test = sig_clf.predict_proba(X_test)
    predict_y_train = sig_clf.predict_proba(X_train)
    predict_test.append(clf.predict(X_test))
    predict_train.append(clf.predict(X_train))
    log_error_test_array.append(log_loss(y_test, predict_y_test, labels=clf.classes_, eps=1e-15))
    log_error_train_array.append(log_loss(y_train, predict_y_train, labels=clf.classes_, eps=1e-15))
    print('When alpha is {} log loss is {}'.format(i,log_loss(y_test, predict_y_test, labels=clf.classes_, eps=1e-15)))

fig, ax = plt.subplots()
ax.plot(alpha, log_error_test_array,c='g')
for i, txt in enumerate(np.round(log_error_test_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_test_array[i]))
plt.grid(linestyle = '-')
plt.title("Cross Validation Error")
plt.xlabel("Alpha")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(log_error_test_array)

print('For values of best alpha {} the train log loss is {}'.format(alpha[best_alpha],log_error_train_array[best_alpha]))

print('For values of best alpha {} the test log loss is {}'.format(alpha[best_alpha],log_error_test_array[best_alpha]))
```

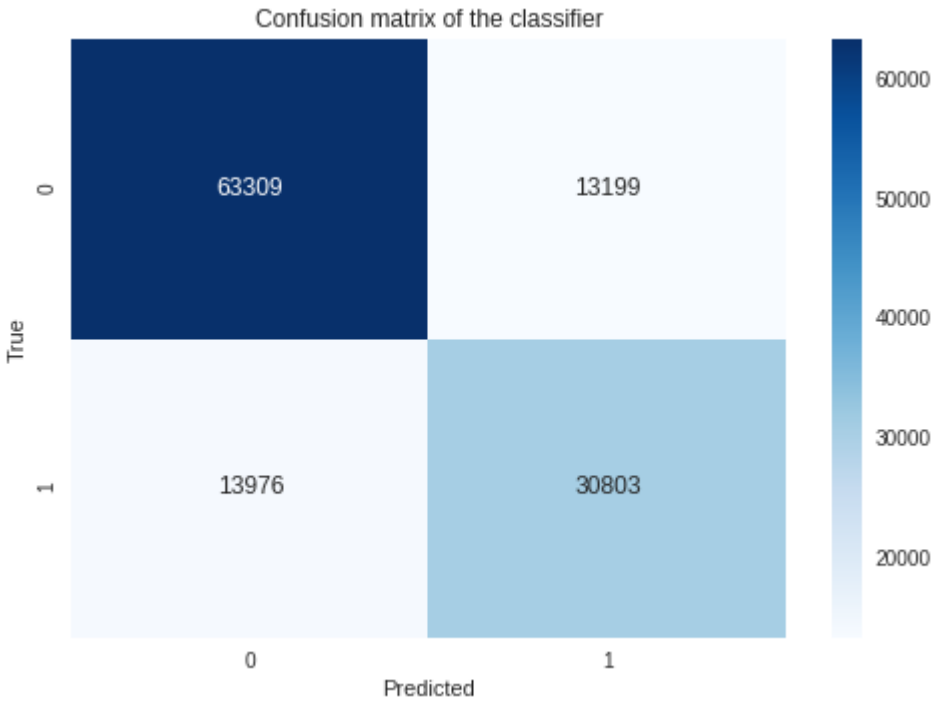
When alpha is 0.0001 log loss is 0.44304936605114953  
When alpha is 0.001 log loss is 0.44207942163087655  
When alpha is 0.01 log loss is 0.4566526366972352  
When alpha is 0.1 log loss is 0.4553373139363503  
When alpha is 1 log loss is 0.48381342523053994  
When alpha is 10 log loss is 0.5454418921782209  
When alpha is 100 log loss is 0.5841136291131219  
When alpha is 1000 log loss is 0.6526084011103247



For values of best alpha 0.001 the train log loss is 0.442749096808023  
For values of best alpha 0.001 the test log loss is 0.44207942163087655



```
In [0]: from pandas_ml import ConfusionMatrix
y_true = np.array(y_test)
y_pred = predict_test[best_alpha]
#print(confusion_matrix(y_test, y_pred))
cm = ConfusionMatrix(y_test,y_pred) #This the confusion matrix of pandas_ml which provides interesting s
confusion_matrix_plot = confusion_matrix(y_test,y_pred) #We are plotting confusion matrix of sklearn
heatmap = sns.heatmap(confusion_matrix_plot, annot=True,cmap='Blues', fmt='g')
plt.title('Confusion matrix of the classifier')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
print("***50)
print("The True Positive Rate observed is",cm.TPR) #This prints the True Positive Rate of the confusion matrix (using pandas_ml confusion matrix).
print("The True Negative Rate observed is",cm.TNR)
print("The False Positive Rate observed is",cm.FPR)
print("The False Negative Rate observed is",cm.FNR)
print("***50)
print("The stats observed for confusion matrix are:")
cm.print_stats()#Prints all the stats of the confusion matrix plotted (using pandas_ml confusion matrix)
```



\*\*\*\*\*

The True Positive Rate observed is 0.3631870158613058  
The True Negative Rate observed is 0.6399841577187115  
The False Positive Rate observed is 0.3600158422812885  
The False Negative Rate observed is 0.6368129841386942

\*\*\*\*\*

The stats observed for confusion matrix are:  
population: 36279  
P: 13555  
N: 22724  
PositiveTest: 13104  
NegativeTest: 23175  
TP: 4923  
TN: 14543  
FP: 8181  
FN: 8632  
TPR: 0.3631870158613058  
TNR: 0.6399841577187115  
PPV: 0.3756868131868132  
NPV: 0.627529665587918  
FPR: 0.3600158422812885  
FDR: 0.6243131868131868  
FNR: 0.6368129841386942  
ACC: 0.5365638523663827  
F1\_score: 0.36933118271503057  
MCC: 0.0031937458432245293  
informedness: 0.0031711735800172836  
markedness: 0.0032164787747310797  
prevalence: 0.37363212877973484  
LRP: 1.0088084278734033  
LRN: 0.9950449186252965  
DOR: 1.0138320481723797  
FOR: 0.372470334412082

## Text Featurization II

### TFIDF Word2Vec

```
In [0]: # avoid decoding problems
df = pd.read_csv("train.csv")

# encode questions to unicode
# https://stackoverflow.com/a/6812069
# ----- python 2 -----
# df['question1'] = df['question1'].apply(lambda x: unicode(str(x),"utf-8"))
# df['question2'] = df['question2'].apply(lambda x: unicode(str(x),"utf-8"))
# ----- python 3 -----
df['question1'] = df['question1'].apply(lambda x: str(x))
df['question2'] = df['question2'].apply(lambda x: str(x))
```

```
In [0]: df.head()
```

Out[3]:

	id	qid1	qid2	question1	question2	is_duplicate
0	0	1	2	What is the step by step guide to invest in sh...	What is the step by step guide to invest in sh...	0
1	1	3	4	What is the story of Kohinoor (Koh-i-Noor) Dia...	What would happen if the Indian government sto...	0
2	2	5	6	How can I increase the speed of my internet co...	How can Internet speed be increased by hacking...	0
3	3	7	8	Why am I mentally very lonely? How can I solve...	Find the remainder when $23^{24}$ is divided by 29	0
4	4	9	10	Which one dissolve in water quickly sugar, salt...	Which fish would survive in salt water?	0

```
In [0]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
# merge texts
questions = list(df['question1']) + list(df['question2'])

tfidf = TfidfVectorizer(lowercase=False, )
tfidf.fit_transform(questions)

# dict key:word and value:tf-idf score
word2tfidf = dict(zip(tfidf.get_feature_names(), tfidf.idf_))
```

- After we find TF-IDF scores, we convert each question to a weighted average of word2vec vectors by these scores.
- here we use a pre-trained GLOVE model which comes free with "Spacy". <https://spacy.io/usage/vectors-similarity> (<https://spacy.io/usage/vectors-similarity>)
- It is trained on Wikipedia and therefore, it is stronger in terms of word semantics.



```
vecs1 = []
# https://github.com/noamraph/tqdm
# tqdm is used to print the progress bar
for qu1 in tqdm(list(df['question1'])):
    doc1 = nlp(qu1)
    # 384 is the number of dimensions of vectors
    mean_vec1 = np.zeros([len(doc1), 384])
    for word1 in doc1:
        # word2vec
        vec1 = word1.vector
        # fetch df score
        try:
            idf = word2tfidf[str(word1)]
        except:
            idf = 0
        # compute final vec
        mean_vec1 += vec1 * idf
    mean_vec1 = mean_vec1.mean(axis=0)
    vecs1.append(mean_vec1)
df['q1_feats_m'] = list(vecs1)
```

```
100%|███████████████████████████████████████████████████████████████████| 404290/404290 [2:13:51<00:00, 50.34it/s]
```

[illegible]

```
df1 = dfnlp.drop(['qid1', 'qid2', 'question1', 'question2'], axis=1)
df2 = dfppro.drop(['qid1', 'qid2', 'question1', 'question2', 'is_duplicate'], axis=1)
df3 = df.drop(['qid1', 'qid2', 'question1', 'question2', 'is_duplicate'], axis=1)
df3_q1 = pd.DataFrame(df3.q1_feats_m.values.tolist(), index= df3.index)
df3_q2 = pd.DataFrame(df3.q2_feats_m.values.tolist(), index= df3.index)
```

id	is_duplicate	cwc_min	cwc_max	csc_min	csc_max	ctc_min	ctc_max	last_word_eq	first_word_eq	abs_len_diff	mean_len	token_set_ratio	token_sort_ratio	fuzz_ratio	fuzz_partial_ratio	longest_substr_ratio	
0	0	0	0.999980	0.833319	0.999983	0.999983	0.916659	0.785709	0.0	1.0	2.0	13.0	100	93	93	100	0.982759
1	1	0	0.799984	0.399996	0.749981	0.599988	0.699993	0.466664	0.0	1.0	5.0	12.5	86	63	66	75	0.596154
2	2	0	0.399992	0.333328	0.399992	0.249997	0.399996	0.285712	0.0	1.0	4.0	12.0	66	66	54	54	0.166667
3	3	0	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.0	0.0	2.0	12.0	36	36	35	40	0.039216
4	4	0	0.399992	0.199998	0.999950	0.666644	0.571420	0.307690	0.0	1.0	6.0	10.0	67	47	46	56	0.175000

	id	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_Total	word_share	freq_q1+q2	freq_q1-q2
0	0	1	1	66	57	14	12	10.0	23.0	0.434783	2	0
1	1	4	1	51	88	8	13	4.0	20.0	0.200000	5	3
2	2	1	1	73	59	14	10	4.0	24.0	0.166667	2	0
3	3	1	1	50	65	11	9	0.0	19.0	0.000000	2	0
4	4	3	1	76	39	13	7	2.0	20.0	0.100000	4	2

	0	1	2	3	4	5	6	7	8	9	...	374	375	376	377	378	379	380	381
0	121.929927	100.083900	72.497894	115.641800	-48.370870	34.619058	-172.057787	-92.502617	113.223315	50.562441	...	12.397642	40.909519	8.150261	-15.170692	18.007709	6.166999	-30.124163	3.700902
1	-78.070939	54.843781	82.738482	98.191872	-51.234859	55.013510	-39.140730	-82.692352	45.161489	-9.556289	...	-21.987077	-12.389279	20.667979	2.202714	-17.142454	-5.880972	-10.123963	-4.890663
2	-5.355015	73.671810	14.376365	104.130241	1.433537	35.229116	-148.519385	-97.124595	41.972195	50.948731	...	3.027700	14.025767	-2.960312	-3.206544	4.355141	2.936152	-20.199555	9.816351
3	5.778359	-34.712038	48.999631	59.699204	40.661263	-41.658731	-36.808594	24.170655	0.235600	-29.407290	...	13.100007	1.405670	-1.891076	-7.882638	18.000561	12.106918	-10.507835	5.243834
4	51.138220	38.587312	123.639488	53.333041	-47.062739	37.356212	-298.722753	-106.421119	106.248914	65.880707	...	13.906532	43.461721	11.519207	-22.468284	45.431128	8.161224	-35.373910	7.728865

5 rows  $\times$  384 columns

	0	1	2	3	4	5	6	7	8	9	...	374	375	376	377	378	379	380	381	
0	125.983301	95.636485	42.114702	95.449980	-37.386295	39.400078	-148.116070	-87.851475	110.371966	62.272814	...	16.165592	33.030668	7.019996	-14.793959	15.437511	8.199658	-25.070834	1.571619	1.603111
1	-106.871904	80.290331	79.066297	59.302092	-42.175328	117.616655	-144.364237	-127.131513	22.962533	25.397575	...	-4.901128	-4.565393	41.520751	-0.727564	-16.413776	-7.373778	2.638877	-7.403457	2.703111
2	7.072875	15.513378	1.846914	85.937583	-33.808811	94.702337	-122.256856	-114.009530	53.922293	60.131814	...	8.359966	-2.165985	10.936580	-16.531660	14.681230	15.633759	-1.210901	14.183826	11.703111
3	39.421531	44.136989	-24.010929	85.265863	-0.339022	-9.323137	-60.499651	-37.044763	49.407848	-23.350150	...	3.311411	3.788879	13.398598	-6.592596	6.437365	5.993293	2.732392	-3.727647	5.613111
4	31.950101	62.854106	1.778164	36.218768	-45.130875	66.674880	-106.342341	-22.901008	59.835938	62.663961	...	-2.403870	11.991204	8.088483	-15.090201	8.375166	1.727225	-6.601129	11.317413	11.544111

5 rows  $\times$  384 columns

```
In [0]: print("Number of features in nlp dataframe :", df1.shape[1])
print("Number of features in preprocessed dataframe :", df2.shape[1])
print("Number of features in question1 w2v dataframe :", df3_q1.shape[1])
print("Number of features in question2 w2v dataframe :", df3_q2.shape[1])
print("Number of features in final dataframe :", df1.shape[1]+df2.shape[1]+df3_q1.shape[1]+df3_q2.shape[1])
```

Number of features in nlp dataframe : 17  
Number of features in preprocessed dataframe : 12  
Number of features in question1 w2v dataframe : 384  
Number of features in question2 w2v dataframe : 384  
Number of features in final dataframe : 794

```
In [0]: # storing the final features to csv file
if not os.path.isfile('final_features.csv'):
    df3_q1['id']=df1['id']
    df3_q2['id']=df1['id']
    df1 = df1.merge(df2, on='id',how='left')
    df2 = df3_q1.merge(df3_q2, on='id',how='left')
    result = df1.merge(df2, on='id',how='left')
    result.to_csv('final_features.csv')
```

```
In [0]: final_features = pd.read_csv('final_features.csv')
final_features.shape
```

Out[6]: (404290, 797)

```
In [0]: y_feats = final_features['is_duplicate']
final_features.drop(['Unnamed: 0', 'id','is_duplicate'], axis=1, inplace=True)
```

```
In [0]: y_feats.shape
```

Out[8]: (404290,)

```
In [0]: final_features.head()
```

Out[8]:

	cwc_min	cwc_max	csc_min	csc_max	ctc_min	ctc_max	last_word_eq	first_word_eq	abs_len_diff	mean_len	...	374_y	375_y	376_y	377_y	378_y	379_y	380_y	381_y	382_y
0	0.999980	0.833319	0.999983	0.999983	0.916659	0.785709	0.0	1.0	2.0	13.0	...	16.165592	33.030668	7.019996	-14.793959	15.437511	8.199658	-25.070834	1.571619	1.603738
1	0.799984	0.399996	0.749981	0.599988	0.699993	0.466664	0.0	1.0	5.0	12.5	...	-4.901128	-4.565393	41.520751	-0.727564	-16.413776	-7.373778	2.638877	-7.403457	2.703070
2	0.399992	0.333328	0.399992	0.249997	0.399996	0.285712	0.0	1.0	4.0	12.0	...	8.359966	-2.165985	10.936580	-16.531660	14.681230	15.633759	-1.210901	14.183826	11.703135
3	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.0	0.0	2.0	12.0	...	3.311411	3.788879	13.398598	-6.592596	6.437365	5.993293	2.732392	-3.727647	5.614115
4	0.399992	0.199998	0.999950	0.666644	0.571420	0.307690	0.0	1.0	6.0	10.0	...	-2.403870	11.991204	8.088483	-15.090201	8.375166	1.727225	-6.601129	11.317413	11.544603

5 rows × 794 columns



```
In [0]: final_features['is_duplicate'] = y_feats.tolist()
p = final_features.groupby('is_duplicate')
final_features = pd.concat([p.get_group(0).sample(75000),p.get_group(1).sample(75000)])
y_feats = final_features['is_duplicate']
#final_features.drop(columns = ['is_duplicate'],inplace=True)
final_features.drop(['is_duplicate'], axis=1, inplace=True)
```

```
In [0]: x_train,x_test, y_train, y_test = train_test_split(final_features, y_feats, stratify=y_feats, test_size=0.3)
```

```
In [0]: y_train.value_counts()
```

Out[11]: 1 52500  
0 52500  
Name: is\_duplicate, dtype: int64

```
In [0]: print("Number of data points in train data :",x_train.shape)
print("Number of data points in test data :",x_test.shape)
```

Number of data points in train data : (105000, 794)  
Number of data points in test data : (45000, 794)

```
In [0]: print("Number of data points in train data :",x_train.shape)
print("Number of data points in test data :",x_test.shape)
```

Number of data points in train data : (105000, 794)  
Number of data points in test data : (45000, 794)

```
In [0]: print("-"*10, "Distribution of output variable in train data", "-"*10)
train_distr = Counter(y_train)
train_len = len(y_train)
print("Class 0: ",int(train_distr[0])/train_len,"Class 1: ", int(train_distr[1])/train_len)
print("-"*10, "Distribution of output variable in train data", "-"*10)
test_distr = Counter(y_test)
test_len = len(y_test)
print("Class 0: ",int(test_distr[1])/test_len, "Class 1: ",int(test_distr[1])/test_len)
```

----- Distribution of output variable in train data -----  
Class 0: 0.5 Class 1: 0.5  
----- Distribution of output variable in train data -----  
Class 0: 0.5 Class 1: 0.5

## Machine Learning Models II

### XGBoost

```
In [0]: import xgboost as xgb
from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import ShuffleSplit, RandomizedSearchCV
from sklearn.metrics import accuracy_score,precision_score,recall_score,f1_score,log_loss,hamming_loss,classification_report
import joblib

#my_cv = ShuffleSplit(n_splits = 5).split(x_train)
param_grid = {'n_estimators' : list(range(100,301,100)), 'max_depth' : list(range(1,6,2)), 'learning_rate' : [0.01, 0.1]} #Parameters for grid search
model = xgb.XGBClassifier()
rsearch_log = RandomizedSearchCV(estimator = model, param_distributions = param_grid, scoring = 'neg_log_loss', cv=3)
rsearch_log.fit(x_train,y_train)

print("The optimal base learners for XGBoost using RandomizedSearchCV is",rsearch_log.best_params_['n_estimators'])
print("The optimal max depth for XGBoost using RandomizedSearchCV is",rsearch_log.best_params_['max_depth'])
print("The optimal learning rate for XGBoost using RandomizedSearchCV is",rsearch_log.best_params_['learning_rate'])
print('*'*50)

rpred = rsearch_log.predict(x_test)
rpred_train = rsearch_log.predict(x_train)

sig_clf = CalibratedClassifierCV(rsearch_log, method="sigmoid") #To get probability of each class
sig_clf.fit(x_train, y_train)

predict_y_train = sig_clf.predict_proba(x_train)
predict_y_test = sig_clf.predict_proba(x_test)

print('\nThe test log loss of XGBoost is {}'.format(log_loss(y_test,predict_y_test)))
print('\nThe train log loss of XGBoost is {}'.format(log_loss(y_train,predict_y_train)))

print('\nThe test accuracy of XGBoost is {}'.format(accuracy_score(y_test, rpred) * 100))
print('\nThe train accuracy of XGBoost is {}'.format(accuracy_score(y_train, rpred_train) * 100))
print('\nThe test precision of XGBoost is {}'.format(precision_score(y_test, rpred)*100))
print('\nThe train precision of XGBoost is {}'.format(precision_score(y_train, rpred_train)*100))
print('\nThe test recall of XGBoost is {}'.format(recall_score(y_test, rpred)*100))
print('\nThe train recall of XGBoost is {}'.format(recall_score(y_train, rpred_train)*100))
print('\nThe test f1 score of XGBoost is {}'.format(f1_score(y_test, rpred)*100))
print('\nThe train f1_score of XGBoost is {}'.format(f1_score(y_train, rpred_train)*100))
print('\nThe test hamming of XGBoost is {}'.format(hamming_loss(y_test, rpred)))
print('\nThe train hamming of XGBoost is {}'.format(hamming_loss(y_train, rpred_train)))
print('*'*50)
print(classification_report(y_test, rpred))
print(""*50)
print(classification_report(y_train, rpred_train))
```

The optimal base learners for XGBoost using RandomizedSearchCV is 300  
The optimal max depth for XGBoost using RandomizedSearchCV is 5  
The optimal learning rate for XGBoost using RandomizedSearchCV is 0.1  
\*\*\*\*\*

The test log loss of XGBoost is 0.349438718329845

The train log loss of XGBoost is 0.2601238790231136

The test accuracy of XGBoost is 83.40444444444445%

The train accuracy of XGBoost is 89.32476190476191%

The test precision of XGBoost is 81.32188698116353%

The train precision of XGBoost is 87.09016761583098%

The test recall of XGBoost is 86.72888888888889%

The train recall of XGBoost is 92.33714285714287%

The test f1 score of XGBoost is 83.93840330350997%

The train f1\_score of XGBoost is 89.6369368453168%

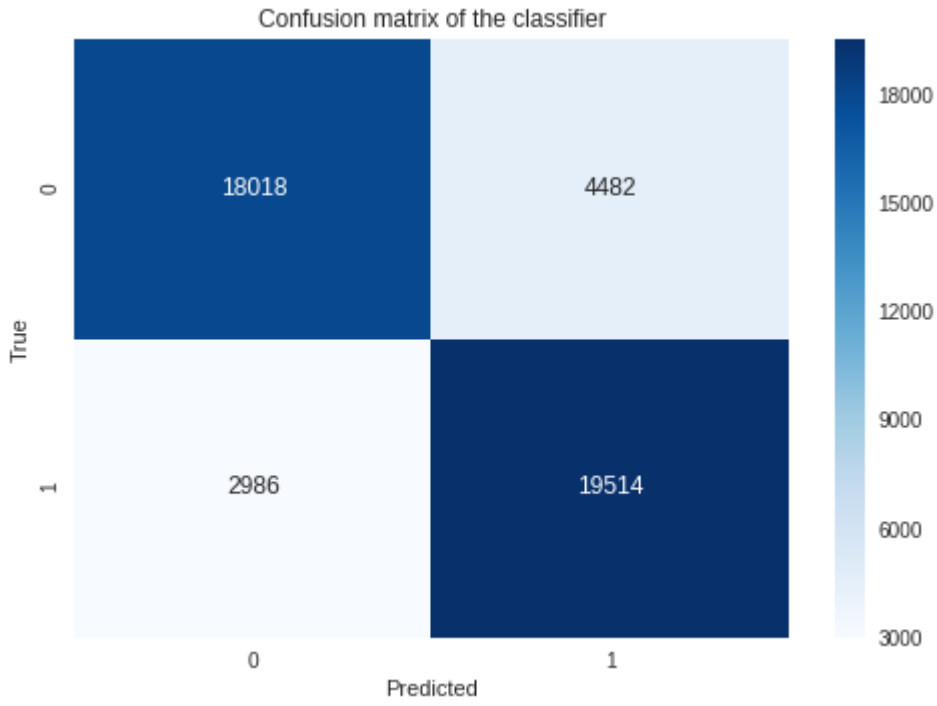
The test hamming of XGBoost is 0.16595555555555555

The train hamming of XGBoost is 0.10675238095238095

*****				
	precision	recall	f1-score	support
0	0.86	0.80	0.83	22500
1	0.81	0.87	0.84	22500
micro avg	0.83	0.83	0.83	45000
macro avg	0.84	0.83	0.83	45000
weighted avg	0.84	0.83	0.83	45000

*****				
	precision	recall	f1-score	support
0	0.92	0.86	0.89	52500
1	0.87	0.92	0.90	52500
micro avg	0.89	0.89	0.89	105000
macro avg	0.89	0.89	0.89	105000
weighted avg	0.89	0.89	0.89	105000

```
In [0]: from pandas_ml import ConfusionMatrix
y_true = np.array(y_test)
y_pred = np.array(rpred)
#print(confusion_matrix(y_test, y_pred))
cm = ConfusionMatrix(y_test,y_pred) #This the confusion matrix of pandas_ml which provides interesting s
confusion_matrix_plot = confusion_matrix(y_test,y_pred) #We are plotting confusion matrix of sklearn
heatmap = sns.heatmap(confusion_matrix_plot, annot=True,cmap='Blues', fmt='g')
plt.title('Confusion matrix of the classifier')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
print("***50)
print("The True Positive Rate observed is",cm.TPR) #This prints the True Positive Rate of the confusion matrix (using pandas_ml confusion matrix).
print("The True Negative Rate observed is",cm.TNR)
print("The False Positive Rate observed is",cm.FPR)
print("The False Negative Rate observed is",cm.FNR)
print("***50)
print("The stats observed for confusion matrix are:")
cm.print_stats()#Prints all the stats of the confusion matrix plotted (using pandas_ml confusion matrix)
```



\*\*\*\*\*
The True Positive Rate observed is 0.5351157222665602
The True Negative Rate observed is 0.47107438016528924
The False Positive Rate observed is 0.5289256198347108
The False Negative Rate observed is 0.46488427773343977
\*\*\*\*\*
The stats observed for confusion matrix are:
population: 5047
P: 2506
N: 2541
PositiveTest: 2685
NegativeTest: 2362
TP: 1341
TN: 1197
FP: 1344
FN: 1165
TPR: 0.5351157222665602
TNR: 0.47107438016528924
PPV: 0.4994413407821229
NPV: 0.5067739204064352
FPR: 0.5289256198347108
FDR: 0.5005586592178771
FNR: 0.46488427773343977
ACC: 0.5028729938577373
F1\_score: 0.5166634559815064
MCC: 0.006202669054356645
informedness: 0.006190102431849365
markedness: 0.0062152611885579745
prevalence: 0.49653259361997226
LRP: 1.0117031624102155
LRN: 0.9868596071183546
DOR: 1.0251743562231759
FOR: 0.4932260795935648

## Summary and Results

- 1) Exploratory data analysis is performed and null values are filled up.
- 2) After getting a sense of data we perform basic feature extraction such as word length, common words, frequency of words etc..
- 3) Later advanced feature extraction such as fuzz ratio, token ratio etc..
- 4) After feature engineering we perform text featurization like TFIDF and apply Logistic Regression and Linear SVM on it.
- 5) The performance seems okay but it can be improved.
- 6) Another text featurization technique TFIDF weighted Word2vec is performed and XGBoost is applied on it.
- 7) XGBoost obtained the best performance so far among all models.

```
In [0]: from prettytable import PrettyTable

x = PrettyTable()
x.field_names = ['Model', 'Vectorization', 'Hyperparameter', 'Train Log Loss', 'Test Log Loss']
x.add_row(['Logistic Regression', 'TFIDF', 'alpha = 0.01', 0.4409, 0.4423])
x.add_row(['', '', '', '', ''])
x.add_row(['Linear SVM', 'TFIDF', 'alpha = 0.001', 0.4427, 0.4420])
x.add_row(['', '', '', '', ''])
x.add_row(['XGBOOST', 'TFIDF weighted word2vec', 'n_estimators = 300', 0.2601, 0.3494])
x.add_row(['', '', 'max_depth = 5', '', ''])
x.add_row(['', '', 'learning_rate = 0.1', '', ''])
print(x.get_string())
```

Model	Vectorization	Hyperparameter	Train Log Loss	Test Log Loss
Logistic Regression	TFIDF	alpha = 0.01	0.4409	0.4423
Linear SVM	TFIDF	alpha = 0.001	0.4427	0.442
XGBOOST	TFIDF weighted word2vec	n_estimators = 300 max_depth = 5 learning_rate = 0.1	0.2601	0.3494